

Code:

```
import numpy as np
import matplotlib.pyplot as plt

# Load dataset
x_train, x_test, y_train, y_test = np.load('regression_data.npy', allow_pickle=True)

# Reshape targets
y_train = y_train.reshape(-1,)
y_test = y_test.reshape(-1,)

# Add bias term (column of ones)
train_data = np.hstack((x_train, np.ones((x_train.shape[0], 1))))
test_data = np.hstack((x_test, np.ones((x_test.shape[0], 1))))

# =====
# Task 1: Standard Linear Regression (Gradient Descent)
# =====
def linear_regression_train(x_train, y_train, lr=1e-3, iterations=7000):
    weight = np.random.randn(2)
    loss = np.zeros(iterations)

    for i in range(iterations):
        y_pred = x_train @ weight # Compute predicted values

        loss[i] = np.mean((y_pred - y_train) ** 2) # Compute MSE

        # Compute gradients
        m_gradient = (2 / x_train.shape[0]) * np.sum((y_pred - y_train) * x_train[:, 0]) # Compute
        gradient for weight
        b_gradient = (2 / x_train.shape[0]) * np.sum(y_pred - y_train)
        # Compute gradient for bias

        # Update weights
        weight[0] -= lr * m_gradient # Apply gradient descent for weight
        weight[1] -= lr * b_gradient # Apply gradient descent for bias

    return weight, loss

weight_standard, loss_standard = linear_regression_train(train_data, y_train)

print("\n=== Standard Linear Regression Parameters ===")
print(f'Weight (m): {weight_standard[0]}') # Print weight[0]
print(f'Bias (b): {weight_standard[1]}') # Print weight[1]
```

```

# =====
# Task 2: Compute MSE
# =====
def compute_mse(y_true, y_pred):
    return np.mean((y_true - y_pred) ** 2)
# Compute MSE formula

y_pred_standard = train_data @ weight_standard.reshape(-1, 1)
y_pred_standard = y_pred_standard.flatten()

print("\n=== Mean Squared Error (Standard Regression) ===")
print(f'MSE: {compute_mse(y_train, y_pred_standard)}')

# =====
# Task 3: Ridge Regression (Gradient Descent)
# =====
def ridge_regression_train(x_train, y_train, lr=1e-3, iterations=7000, lambda_reg=0.1):
    weight = np.random.randn(2)
    loss = np.zeros(iterations)

    for i in range(iterations):
        y_pred = x_train @ weight # Compute predicted values

        loss[i] = np.mean((y_pred - y_train) ** 2) + lambda_reg * np.sum(weight ** 2)
    # Compute MSE with regularization term

    # Compute gradients with regularization
    m_gradient = (2 / x_train.shape[0]) * np.sum((y_pred - y_train) * x_train[:, 0]) + 2 * lambda_reg * weight[0] # Compute weight gradient with regularization
    b_gradient = (2 / x_train.shape[0]) * np.sum(y_pred - y_train) + 2 * lambda_reg * weight[1] # Compute bias gradient

    # Update weights
    weight[0] -= lr * m_gradient # Apply gradient descent for weight
    weight[1] -= lr * b_gradient
    # Apply gradient descent for bias

    return weight, loss

weight_ridge, loss_ridge = ridge_regression_train(train_data, y_train)

print("\n=== Ridge Regression Parameters ===")
print(f'Weight (m): {weight_ridge[0]}')
print(f'Bias (b): {weight_ridge[1]}')

y_pred_ridge = test_data @ weight_ridge.reshape(-1, 1) # Compute predictions for test data
y_pred_ridge = y_pred_ridge.flatten()

```

```

mse_ridge = compute_mse(y_test, y_pred_ridge)

print("\n=== Mean Squared Error (Ridge Regression) ===")
print(f'MSE: {mse_ridge}')

# =====
# Task 4: Plot Loss Curve
# =====
plt.plot(loss_standard, label='Standard Regression')
plt.plot(loss_ridge, label='Ridge Regression')
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.title('Loss Curve Comparison')
plt.legend()
plt.show()

# =====
# Task 5: Closed-form Ridge Regression
# =====
def closed_form_ridge(x_train, y_train, lambda_reg=0.1):
    I = np.eye(x_train.shape[1])
    I[-1, -1] = 0
    w_closed_form = np.linalg.inv(x_train.T @ x_train + lambda_reg * I) @ x_train.T @ y_train #
    Compute closed-form solution (Equation 4.27)
    return w_closed_form

weight_closed_form = closed_form_ridge(train_data, y_train)
y_pred_closed_form = test_data @ weight_closed_form.reshape(-1, 1)
y_pred_closed_form = y_pred_closed_form.flatten()
mse_closed_form = compute_mse(y_test, y_pred_closed_form)

print("\n=== Closed-form Ridge Regression Parameters ===")
print(f'Weight (m): {weight_closed_form[0]}')
print(f'Bias (b): {weight_closed_form[1:]}')
print("\n=== Mean Squared Error (Closed-form Ridge Regression) ===")
print(f'MSE: {mse_closed_form}')

# =====
# Task 6: Predictive Distribution
# =====
# Compute posterior covariance
lambda_reg=0.1
I = np.eye(x_train.shape[1])
sigma_n = lambda_reg * np.eye(x_train.shape[1])

# Compute posterior covariance
S_N_inv = x_train.T @ x_train + sigma_n
S_N = np.linalg.inv(S_N_inv)

```

```

# Compute posterior mean
w_mean = S_N @ x_train.T @ y_train

# Predictive mean
predictive_mean = x_test @ w_mean

# Predictive variance
predictive_variance = np.sum(x_test @ S_N * x_test, axis=1)


# =====
# Task 7: Plot Predictions
# =====
plt.figure(figsize=(10, 6))

# Prediction lines

plt.scatter(x_test[:, 0], y_test, label='Ground Truth', color='blue', alpha=0.6)

# Standard Regression
y_pred_standard = test_data @ weight_standard
plt.scatter(x_test[:, 0], y_pred_standard, label='Standard Regression', color='red')

# Ridge Regression
y_pred_ridge = test_data @ weight_ridge
plt.scatter(x_test[:, 0], y_pred_ridge, label='Ridge Regression', color='green')

# Closed-form Ridge Regression
y_pred_closed_form = test_data @ weight_closed_form
plt.scatter(x_test[:, 0], y_pred_closed_form, label='Closed-form Ridge', color='purple')
plt.xlabel('X')
plt.ylabel('Y')
plt.title('Prediction Results')
plt.legend()
plt.show()

# Confidence Interval Shading

# =====
# Plot Confidence Intervals
# =====
plt.scatter(x_test[:, 0], y_test, label='Ground Truth', color='blue', alpha=0.6)

# Ridge Regression

```

```
plt.scatter(x_test[:, 0], predictive_mean, label='Ridge Regression', color='green')
```

```
plt.xlabel('X')  
plt.ylabel('Y')  
plt.title('Predictive Distribution with Confidence Interval')  
plt.legend()  
plt.show()
```

Output result:

```
=== Standard Linear Regression Parameters ===  
Weight (m): 52.74342253045384  
Bias (b): -0.33378000891502935  
  
=== Mean Squared Error (Standard Regression) ===  
MSE: 99.34237310127232  
  
=== Ridge Regression Parameters ===  
Weight (m): 47.62796542012239  
Bias (b): -0.3949129745921893  
  
=== Mean Squared Error (Ridge Regression) ===  
MSE: 122.25191150505415  
  
=== Closed-form Ridge Regression Parameters ===  
Weight (m): 52.73598698626496  
Bias (b): [-0.3339075]  
  
=== Mean Squared Error (Closed-form Ridge Regression) ===  
MSE: 110.41535937243981
```



