



# Machine Learning

## LABORATORY: Gradient Descent In Class

NAME:

STUDENT ID#:

### Objectives:

- Understand and implement Stochastic Gradient Descent (SGD) using one sample at a time (Algorithm 7.1).
- Apply the sigmoid activation function to perform binary classification.
- Learn how to update model weights manually using the gradient of the loss function.
- Practice using NumPy for mathematical operations on real MNIST data.
- Evaluate classification results using accuracy and misclassified samples visualization.

### Part 1. Instruction

- In this assignment, you will implement Stochastic Gradient Descent (SGD) from scratch using Algorithm 7.1 provided in class.
- The task is to build a binary classifier to distinguish whether an MNIST image corresponds to a target digit or not (e.g., "Is this a 0 or not?").
- You will complete the code template provided in the in-class assignment.
- The classifier will use the sigmoid function for binary probability output. Use only NumPy for all computations. Do not use libraries like scikit-learn or PyTorch.
- Evaluate your results using the printed accuracy and a visualization of misclassified samples.

### Part 2. Arithmetic Instructions.

#### Algorithm 7.1: Stochastic gradient descent

**Input:** Training set of data points indexed by  $n \in \{1, \dots, N\}$

Error function per data point  $E_n(\mathbf{w})$

Learning rate parameter  $\eta$

Initial weight vector  $\mathbf{w}$

**Output:** Final weight vector  $\mathbf{w}$

$n \leftarrow 1$

**repeat**

$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla E_n(\mathbf{w})$  // update weight vector

$n \leftarrow n + 1 \pmod{N}$  // iterate over data

**until** convergence

**return**  $\mathbf{w}$



### Part 3. Code Template.

Step	Procedure
1	<pre> #Load Dataset import struct import numpy as np import matplotlib.pyplot as plt  # =====Load IDX Files ===== def load_images(filename):     with open(filename, 'rb') as f:         _, num, rows, cols = struct.unpack("&gt;IIII",         f.read(16))         images=np.frombuffer(f.read(),         dtype=np.uint8)         images = images[: (len(images) // (rows * cols))         * rows * cols]         return images.reshape(-1, rows *         cols).astype(np.float32) / 255.0  def load_labels(filename):     with open(filename, 'rb') as f:         _, num = struct.unpack("&gt;II", f.read(8))         labels = np.frombuffer(f.read(),         dtype=np.uint8)         return labels[:num] </pre>
2	<pre> # ===== 1. Sigmoid Function ===== def sigmoid(z):     # TODO: Implement sigmoid function     pass  # ===== 2. SGD: Algorithm 7.1 ===== def sgd_logistic(X, y, eta, max_iters):      pass </pre>
3	<pre> # =====Show Misclassified Samples ===== def show_misclassified(X, true_labels, pred_labels, max_show=10):     mis_idx = np.where(true_labels !=     pred_labels)[0][:max_show]     plt.figure(figsize=(10, 2))     for i, idx in enumerate(mis_idx):         plt.subplot(1, len(mis_idx), i + 1)         plt.imshow(X[idx, 1:].reshape(28, 28), cmap='gray')         plt.axis('off')         plt.title(f"T:{true_labels[idx]} P:{pred_labels[idx]}")     plt.suptitle("Misclassified Samples")     plt.show() </pre>
4	<pre> # ===== 3. Main ===== if __name__ == "__main__":     # == Load Data ==     X_train = load_images("train-images.idx3-ubyte")     y_train = load_labels("train-labels.idx1-ubyte")     X_test = load_images("t10k-images.idx3-ubyte")     y_test = load_labels("t10k-labels.idx1-ubyte") </pre>



```

# === Choose binary classification target digit ===
TARGET_DIGIT = 0 # TODO: Fill in (0 to 9)

y_train_bin = np.where(y_train == TARGET_DIGIT, 1, 0)
y_test_bin = np.where(y_test == TARGET_DIGIT, 1, 0)

# === Add bias term ===
X_train = np.hstack([np.ones((X_train.shape[0], 1)),
X_train])
X_test = np.hstack([np.ones((X_test.shape[0], 1)),
X_test])

# === Set parameters ===
#eta = _____ # TODO: Learning rate
#max_iters = _____ # TODO: Number of SGD iterations

# === Train ===
#w = sgd_logistic(X_train, y_train_bin, eta, max_iters)

# === Predict ===
#pred_probs = _____
#preds = (pred_probs >= 0.5).astype(int)

# === Evaluate ===
#accuracy = np.mean(preds == y_test_bin)
#print(f"Accuracy: {accuracy:.4f}")

# === Show Misclassified Samples ===
#show_misclassified(X_test, y_test_bin, preds)

```

5

#### #Example Output:

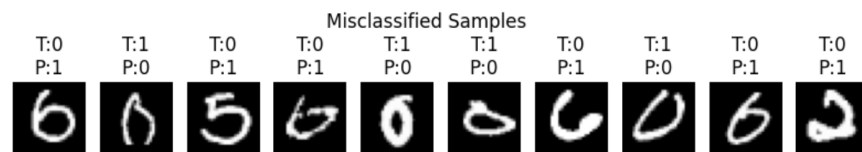
```

[INFO] Header: 60000 images, 28x28
[INFO] Loading 60000 images based on file size
[INFO] Loading 60000 labels based on file size
[INFO] Header: 10000 images, 28x28
[INFO] Loading 10000 images based on file size
[INFO] Loading 10000 labels based on file size

```

```
[INFO] Binary classification: '0' vs not-0
```

```
Test Accuracy (is 0 or not): 0.9927
```



## Grading Assignment & Submission (30% Max)

### Implementation:

1. (15%) Implement the Stochastic Gradient Descent (SGD) algorithm based on Algorithm 7.1.
2. (10%) The model runs successfully without errors, Trains using the provided MNIST dataset, and output the test accuracy.
3. (5%) Set the class of binary classification to the last digit of your student ID. (e.g., if your ID ends in 7, use the class '7'). Displays at least 5 misclassified test images along with their true (T:) and predicted (P:) labels, as shown in the "Example Output" in the last pages of this document.



## Submission :


1. Report: Answer all conceptual questions. Include screenshots of your results in the last pages of this PDF File.
2. Code: Submit your complete Python script in either .py or .ipynb format.
3. Upload both your report and code to the E3 system (**Labs4 In Class Assignment**). Name your files correctly:
  - a. Report: StudentID\_Lab4\_InClass.pdf
  - b. Code: StudentID\_Lab4\_InClass.py or StudentID\_Lab4\_InClass.ipynb
4. Deadline: 16:20 PM
5. Plagiarism is **strictly prohibited**. Submitting copied work from other students will result in penalties.

## Example Output (Just for reference):

```
[INFO] Header: 60000 images, 28x28
[INFO] Loading 60000 images based on file size
[INFO] Loading 60000 labels based on file size
[INFO] Header: 10000 images, 28x28
[INFO] Loading 10000 images based on file size
[INFO] Loading 10000 labels based on file size
```

```
[INFO] Binary classification: '0' vs not-0
```

```
Test Accuracy (is 0 or not): 0.9927
```



## Code Results and Answer:

