



Machine Learning

LABORATORY: CNN In Class

NAME:

STUDENT ID#:

Objectives:

- Students will implement 2D convolution and pooling operations from scratch (no PyTorch/TensorFlow high-level API like `nn.Conv2d`, `F.max_pool2d`, etc.) using only NumPy.
- Understand padding and stride behavior. Apply vertical and horizontal edge detection. Visualize results in black & white

Part 1. Instruction

- In this assignment, you will implement a basic Convolutional Neural Network operation pipeline using NumPy only — without using any deep learning libraries like PyTorch, TensorFlow, or OpenCV's built-in convolution functions.
- You will manually implement a general 2D convolution function that supports:
 - Padding (to preserve spatial dimensions)
 - Stride (to downsample the output)
- Then, you will apply vertical and horizontal edge detection filters to a grayscale input image and visualize the effects of:
 - Padding (`padding=1`)
 - Strided convolution (`stride=2`)
- Specifically, your tasks are to:
 - Load and normalize a grayscale image (e.g., `checkerboard.png`) for testing edge detection.
 - Implement the general convolution operation

$$C(j, k) = \sum_l \sum_m I(j + l, k + m) K(l, m)$$

$I(j + l, k + m)$ = **Image region**

$K(l, m)$ = **Kernel (Filter)**

$C(j, k)$ = **output at position (j, k)**

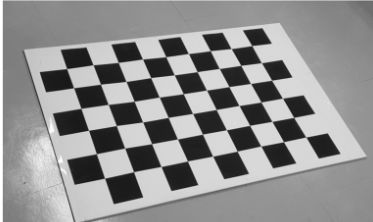
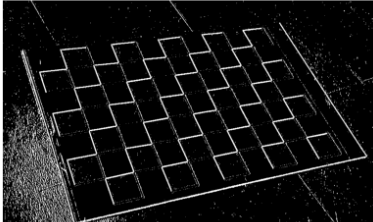
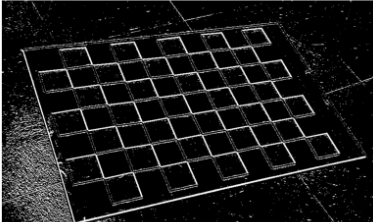
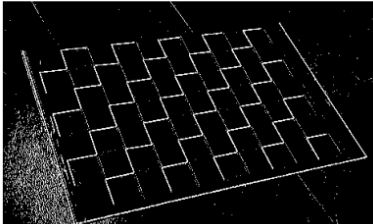
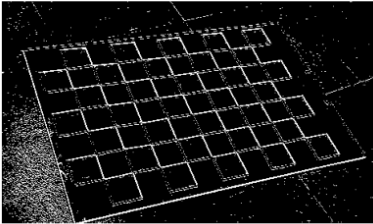
- Apply the vertical edge detection kernel and horizontal edge detection kernel from Slide 6 to detect pattern structures in the image.
 - Apply convolution again using `stride=2` to observe how spatial resolution changes (see Slide 8).
 - Visualize the output as black-and-white (binary) images using thresholding.
- At the end of the lab, please answer the two short questions to demonstrate your understanding of padding and stride.



Part 2. Code Template

Step	Procedure
1	<pre>import numpy as np import cv2 import matplotlib.pyplot as plt # Step 1: Load a grayscale image and normalize # ► Slide 5: Understanding image representation image = cv2.imread('original.png', cv2.IMREAD_GRAYSCALE) image = image.astype(np.float32) / 255.0 # Normalize to range [0, 1]</pre>
2	<pre># Step 2: General Convolution Function # ► Slide 8: $C(j, k) = \sum_l \sum_m I(j + l, k + m) * K(l, m)$ def convolve2d(image, kernel, padding=0, stride=1): # TODO 1: Flip kernel for convolution # TODO 2: Apply zero-padding if padding > 0 # TODO 3: Calculate output height and width # TODO 4: Slide the kernel across the image with stride # TODO 5: At each position, compute the sum of element-wise multiplication return np.zeros((1, 1)) # Placeholder, replace with real output</pre>
3	<pre># Step 3: Define edge detection filters # ► Slide 6: Vertical & Horizontal edge filters vertical_filter = np.array([[x, x, x], [x, x, x], [x, x, x]], dtype=np.float32) horizontal_filter = np.array([[x, x, x], [x, x, x], [x, x, x]], dtype=np.float32)</pre>
4	<pre># Step 4: Convolve image with filters (padding=1, stride=1) # ► Slide 7: Padding helps preserve image size # TODO: vertical_edges = convolve2d(image, vertical_filter, padding=1) # TODO: horizontal_edges = convolve2d(image, horizontal_filter, padding=1) # Try strided convolutions (padding=1, stride=2) # ► Slide 8: Stride reduces spatial resolution # TODO: vertical_stride = convolve2d(image, vertical_filter, padding=1, stride=2) # TODO: horizontal_stride = convolve2d(image, horizontal_filter, padding=1, stride=2)</pre>
5	<pre># Step 5: Visualization and Binarization function for black-and-white display def binarize(img, threshold=0.5): img = img - np.min(img) if np.max(img) != 0: img = img / np.max(img) return (img > threshold).astype(np.float32)</pre>



	<pre># Visualization # TODO: Use matplotlib to show: # - Original image # - Vertical edges (pad=1) # - Horizontal edges (pad=1) # - Vertical edges (stride=2) # - Horizontal edges (stride=2)</pre>
5	<p>#Example Output: (References Only)</p> <div> <div>Original</div>  </div> <div> <div>Vertical Edge (pad=1)</div>  </div> <div> <div>Horizontal Edge (pad=1)</div>  </div> <div> <div>Vertical Edge (stride=2)</div>  </div> <div> <div>Horizontal Edge (stride=2)</div>  </div>

Grading Assignment & Submission (30% Max)

Implementation:

- (10%) Correctly implement the `convolve2d()` function, including kernel flipping, padding, and stride
- (5%) Correctly apply vertical and horizontal edge detection filters.
- (5%) Apply binary thresholding to convert outputs to black-and-white images, and clearly visualize them using matplotlib. Show all five views: original, vertical (pad=1), horizontal (pad=1), vertical (stride=2), and horizontal (stride=2).

Question:

- (5%) What types of patterns are detected by the vertical edge filter in an image? How is this different from the horizontal edge filter?
- (5%) What is the effect of padding on the output image when applying convolution? Why is padding used?

Submission :

- Report: Provide your screenshots of your results in the last pages of this PDF File.
- Code: Submit your complete Python script in either .py or .ipynb format.
- Upload both your report and code to the E3 system (**Labs6 In Class Assignment**). Name your files correctly:
 - Report: StudentID_Lab6_InClass.pdf
 - Code: StudentID_Lab6_InClass.py or StudentID_Lab6_InClass.ipynb
- Deadline: 16:20 PM
- Plagiarism is **strictly prohibited**. Submitting copied work from other students will result in penalties.



Results and Discussion:

